

CSC411: Project 2

Due on Sunday, February 18, 2018

Ying Yang

March 20, 2018

Part 1

Dataset description

Describe the datasets.

The dataset consists of headlines of real news and headlines of fake news, each stored in a separate file. All headlines' punctuations and apostrophes are removed so that it only contain clean words.

We have a dataset of 3266 headlines

real news headlines: 1968

fake news headlines: 1298

3 examples of words that may be useful

We collected words that appear frequently in real and not frequently in fake news, and vice versa, since if a word appears frequently in a one and not frequent in the other, then the word is will strongly predict the headline. Among the results, we selected words that are meaningful (we didn't select words that are english grammar words, like "the", "and" , etc.)

examples of words:

word that appears frequent in real, and less frequent in fake: ['criticism', 10, 1]

words that appear frequent in fake, and less frequent in real: ['reporter', 15, 1] ['liberty', 12, 1]

Test scripts of the project:

For the following problems, test scripts are run simply by uncommenting out the part # in the *main()* function.

Listing 1: Test scripts

```
5         if __name__ == "__main__":
10             # part1()
               # part2()
               # part3a()
               # part3b()
               # part4()
               # part5()
               # part6()
               # part7()
               # part8()
```

Part 2

Problem: Implement the Naive Bayes algorithm for predicting whether a headline is real or fake.

Implementation details: We used naive bayes probability to predict whether a headline is real or fake. For each of the headline, we used naive bayes to calculate $P(\text{real} \mid \text{headline})$ and $P(\text{fake} \mid \text{headline})$ and compare which one has greater probability. For example, if $P(\text{real} \mid \text{headline}) > P(\text{fake} \mid \text{headline})$, then we predict it to be real. We check if the prediction is correct using the corresponding label. We count the number of correct predictions for all headlines and divide it by the total number of headlines in each of training_set, validation_set and test_test to get the performance.

In particular, we used the formula:

$$P(\text{real} \mid \text{headline}) = \text{likelihood_probability}(\text{headline}) * \text{prior_probability}(\text{real})$$

$$P(\text{fake} \mid \text{headline}) = \text{likelihood_probability}(\text{headline}) * \text{prior_probability}(\text{fake})$$

$$= P(\text{real} \mid \text{headline}) = \sum \log(P(w_i \mid \text{real})) + \log(\text{prior_probability}(\text{real}))$$

$$= P(\text{fake} \mid \text{headline}) = \sum \log(P(w_i \mid \text{fake})) + \log(\text{prior_probability}(\text{fake}))$$

Step 1 : Calculate the likelihood of a headline.

(similar way for getting likelihood for fake news headline)

Step 1.1 We first calculate the total number of words that are in real and fake news in the dataset.

Step 1.2 We then maintained a table of real and fake words and its counts. (in python dictionary) For each word in the dataset, we store the word as key and $P(\text{word} \mid \text{real})$ and $P(\text{word} \mid \text{fake})$ as a list of value of the word.

$$P(\text{word} \mid c) = (\text{count}(\text{word}, c) + m * p_hat) / (\text{count}(c) + m), c = \text{real or fake}$$

m and p_hat are the parameters that we will tune with validation set to get best performance.

We used validation set to tune the parameters m and p

Table of probability for each word ($P(w \mid \text{real})$, $P(w \mid \text{fake})$)

Listing 2: Make a table of probability for each word

```

probability_dict = {}
for word in probability_dict:
    # get real words probability
    if word in words_occurrences_real_table:
5         real_word_likelihood = float(words_occurrences_real_table[word] + m * p_hat) /
            float(total_real_healines + m)
    else:
        real_word_likelihood = float(m * p_hat) / float(total_real_healines + m)
        probability_dict[word] = [real_word_likelihood]
10    # get fake words probability
    if word in words_occurrences_fake_table:
        fake_word_likelihood = float(words_occurrences_fake_table[word] + m * p_hat) /
            float(total_fake_healines + m)
    else:
15         fake_word_likelihood = float(m * p_hat) / float(total_fake_healines + m)
        probability_dict[word].append(fake_word_likelihood)

```

Step 1.3 Calculate likelihood of the headline. We do so by adding $\log(p(\text{word} \mid c))$ if word appears in the

headline, $\log(1 - p(\text{word} | c))$ if words doesn't appear in the headline, for each of the words in the table. This sum is the likelihood of the headline. We will called this likelihood_headline

Listing 3: Calculate likelihood for a headline

```

likelihood_headline_real = 0
likelihood_headline_fake = 0

    for word in likelihood_table:
        if word in headline:
            likelihood_headline_real += math.log(likelihood_table[word][0])
            likelihood_headline_fake += math.log(likelihood_table[word][1])
        else:
            likelihood_headline_real += math.log(1 - likelihood_table[word][0])
            likelihood_headline_fake += math.log(1 - likelihood_table[word][1])

```

Step 2: Calculate prior probability of real or fake.

$P(\text{real}) = \text{total number of headlines real} / \text{total headlines}$ $P(\text{fake}) = \text{total number of headlines fake} / \text{total headlines}$

Step 3: Calculate the posterior probability for a headline being real using $\text{likelihood_headline} + \log(\text{prior}(\text{real}))$. We do similar operation for calculating posterior probability for a headline being fake.

Step 4: Count the accuracy.

We use the posterior probability for each headline being real or fake and compare which one is greater, and use the one that is greater to predict the whether a headline is real or fake. For example, if a headline's posterior probability for being real is greater than the posterior probability that of being fake, then we predict it is real. We iterate through all the headlines and check with the label of it. We increment correct guesses number if prediction is correct.

Listing 4: Calculate performance for a dataset

```

correct = 0
    for i in range(len(dataset)):
        headline = dataset[i]
        if (predict_headline(likelihood_table, headline, dataset, dataset_label) ==
            "real" and dataset_label[i] == 1) or \
            (predict_headline(likelihood_table, headline, dataset, dataset_label) ==
            "fake" and dataset_label[i] == 0):
            correct += 1

    return (float(correct) / len(dataset)) * 100

```

Step 5: We tune the parameters m and p_{hat} using validation set.

To do so, we tried with an array of numbers

$m = [1.0, 0.5, 1e1, 1e2, 1e3, 1e4]$, $p_{\text{hat}} = [0.9, 3e-1, 1e-1, 3e-2, 1e-2, 3e-3]$

After several tests and changing the values for p_{hat} and m , we found out that $m=0.5$ and $p_{\text{hat}} = 0.1$ are the best parameters.

Results:

training set performance: 97.8127734033%

validation set performance: 99.387755102%

test set performance: 98.0592441267%

Part 3

Part 3a

How we obtained the lists:

We have maintained a table of probabilities for each word ($P(w \mid \text{real})$ and $P(w \mid \text{fake})$). For each of the word, we calculate store its posterior probability by doing:

Presence predict real: $P(\text{real} \mid w) = P(w \mid \text{real}) + \log(\text{prior_prob}(\text{real}))$

Presence predict fake: $P(\text{fake} \mid w) = P(w \mid \text{fake}) + \log(\text{prior_prob}(\text{fake}))$

Absence predict real: $P(\text{real} \mid w) = 1 - P(w \mid \text{fake}) + \log(\text{prior_prob}(\text{real}))$

Absence predict fake: $P(\text{fake} \mid w) = 1 - P(w \mid \text{fake}) + \log(\text{prior_prob}(\text{fake}))$

Then we iterate through the table to find the top 10 highest values.

List the 10 words whose presence most strongly predicts that the news is real: ('travel')

('turnbull')

('ban')

('korea')

('north')

('says')

('us')

('trumps')

('trump')

('donald')

List the 10 words whose absence most strongly predicts that the news is real:

('if')

('are')

('you')

('clinton')

('and')

('just')

('is')

('a')

('hillary')

('the')

List the 10 words whose presence most strongly predicts that the news is fake:

('if')

('are')

('you')

('clinton')

('and')

('just')

('is')

('a')

('hillary')
('the')

List the 10 words whose absence most strongly predicts that the news is fake:

('travel')
('turnbull')
('ban')
('korea')
('north')
('says')
('us')
('trumps')
('trump')
('donald')

Compare the influence of presence vs absence of words on predicting whether the headline is real or fake news.

Part 3b

Write vectorized code that computes the gradient of the cost function with respect to the weights and biases of the network

10 non-stopwords that most strongly predict that the news is real: ('australia') ('wall') ('travel') ('turnbull') ('ban') ('korea') ('north') ('says') ('trumps') ('donald')

10 non-stopwords that most strongly predict that the news is fake: ('black') ('watch') ('win') ('obama') ('new') ('america') ('just') ('clinton') ('hillary') ('trump')

Part 3c

Why might it make sense to remove stop words when interpreting the model?

Because stop words usually doesn't have any meaning. They are only words that constructs syntactical sentences, but doesn't contribute to the meaning of the headline. Since we want to know which are the words that predict strongly, we need only consider words that make sense.

Why might it make sense to keep stop words?

It makes sense when sometimes a fake news might tend to use a set of non-stop english words to make up the title. For example, in real news, one might not use "seems", whereas in fake news one might use the word "seems" a lot when making up a title "that creates false statements", without proofs, such as "Trump seems to have intention to declare war". In this case, real news might not use this word "seems", since they tend to only use solid proofs, and we rarely see "seems" in titles with solid proofs.

Part 4

Train a Logistic Regression model on the same dataset

Implementation details:

Plot the learning curves (performance vs. iteration) of the Logistic Regression model.

Describe how you selected the regularization parameter (and describe the experiments you used to select it).

Part 5

Write vectorized code that performs gradient descent with momentum, and use it to train your network. Plot the learning curves

Listing 5: Gradient descent with momentum

```
train_neural_network(image_prefix, type="None", momentum_term=0):  
  
    v_w = 0 # for momentum  
    v_b = 0  
    if type == "momentum":  
        v_w = momentum_term * v_w + alpha * df_w(train_set.T, W, b, train_label.T)  
        W -= v_w  
        v_b = momentum_term * v_b + alpha * df_b(train_set.T, W, b, train_label.T)  
        b -= v_b  
  
    if __name__ == "__main__":  
        momentum_term = 0.99  
        train_neural_network("part5", "momentum", momentum_term)
```

Part 6

Part 6d

Momentum trajectory walks faster than vanilla gradient descent, and momentum can lead to points past the central local minimum while vanilla gradient descent walks directly to the local point. The cause of these differences is the extra term in momentum which is $\text{momentum_term} * v_w$ speeds up the walking pace, and it retains effects of previous v_m term.

Part 6e

I found proper settings by checking the gradient descent values and trial-and-error. Whichever weights' values are the most among all the gradient descent values, should be chosen as the two variables. It speeds up the descent process. Another setting is momentum term. Choose it by trial-and-error to see which one fits the graph the best. Initial weights values are also important because the directions of original gradient descents can demonstrate the difference between two ways, too.

Other settings do not demonstrate the benefits of momentum are alpha and axis range. Choose plot axis which can zoom in the contour at the proper size and alpha which is neither too large nor too small. Momentum term do not really affect visualization unless it is so large that points past local minimum too much or offtrack too much.

Part 7

Assumption: For each neuron, due to unknown size, suppose matrix multiplication complexity is 1.

For i -th layer, each j in K neurons, backpropagation computes w_j with complexity of $O(K)$ for all i .

If it computes individually, w_j has computation complexity of $O(K^{2i-3})$

Final computation complexity, for backpropagation, is $O(N K^2)$, individually, is $O(K^{2N-2})$

When N is larger than 2, backpropagation is $O(K^{2N-4})$ faster than individual computation on condition that N is known.

Part 8

Use a fully-connected neural network with a single hidden layer, but train using mini-batches using an optimizer

Description of the system:

We initialize inputs and weights as follows:

`dim_x = 64 * 64 * 3`

`dim_h = 500`

`dim_out = 6`

We used one hidden layer of size 500.

The resolution was 64 * 64.

We used Adam as optimizer

The model we used was:

Listing 6: model of the system

```
5      x = Variable(torch.from_numpy(train_set[train_index]),
10                      requires_grad=False).type(dtype_float)
      y_classes = Variable(torch.from_numpy(np.argmax(train_label[train_index], 1)),
                          requires_grad=False).type(dtype_long)

      model = torch.nn.Sequential(
          torch.nn.Linear(dim_x, dim_h),
          torch.nn.ReLU(),
          torch.nn.Linear(dim_h, dim_out),
      )
```

Optimization steps and outputs:

We tried varying parameters, optimizers and functions and other factors that might improve the performance, such as initial weights, learning rate, batch size, number of iterations, optimization function to optimize the training performance, activation function, number of hidden layers, weight of hidden layers.

Listing 7: initial inputs and weights and parameters and functions chosen for optimization

```
5      dim_x = 64 * 64 * 3
      dim_h = 500
      dim_out = 6

      learning_rate = 1e-4
      max_iteration = 500
      optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)
      loss_fn = torch.nn.CrossEntropyLoss()
```

Observations: Activation function: Since we are using mini-batches, we found out that ReLU was the best option. It accelerates the convergence of stochastic gradient descent compared to sigmoid and tanh functions.

Learning rate: We first tried using learning rate of 0.001 then increased to 0.00001. We found out that the performance was best at learning rate = 0.0001. The performance improved from 0.001 to 0.0001 but decreased from 0.0001 to 0.00001.

Number of layers: We experimented with a single and 2 hidden layers, with fixed weight. Keeping other parameters unchanged, we found insignificant difference between having 2 or 1 hidden layers.

Weights: We experimented with weights from 20 to 700. At low weight, the test performance had poor results. The accuracy performance increased slower compared to having setting weight as 700. Also, the test performance was unstable when the weight was low. It might first increase and decrease later.

Learning curve:

Final performance: The performance of the test set was at 85.83%.

The results were:

test performance: 36.6666666667

test performance: 83.3333333333

test performance: 85.0

test performance: 85.0

test performance: 86.6666666667

test performance: 85.8333333333

test performance: 85.8333333333

test performance: 85.8333333333

test performance: 85.8333333333

test performance: 85.8333333333

test performance: 85.8333333333

Part 9

Select two of the actors, and visualize the weights of the hidden units that are useful for classifying input photos as those particular actors. Explain how you selected the hidden units

Actors selected: balwin and bracco

Approach: Try lots of images in a dataset, and see which ones activate the neuron the most.

Part 10

We modified code for AlexNet. We want to extract activation values before it classifies the actor, so we deleted the line `x = x.classifier(x)` in `forward()` method.

The modified code is:

Listing 8: modified code for AlexNet

```

class MyAlexNet(nn.Module):
    def load_weights(self):
        an_builtin = torchvision.models.alexnet(pretrained=True)

        features_weight_i = [0, 3, 6, 8, 10]
        for i in features_weight_i:
            self.features[i].weight = an_builtin.features[i].weight
            self.features[i].bias = an_builtin.features[i].bias

        classifier_weight_i = [1, 4, 6]
        for i in classifier_weight_i:
            self.classifier[i].weight = an_builtin.classifier[i].weight
            self.classifier[i].bias = an_builtin.classifier[i].bias

    def __init__(self, num_classes=1000):
        super(MyAlexNet, self).__init__()
        self.features = nn.Sequential(
            nn.Conv2d(3, 64, kernel_size=11, stride=4, padding=2),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=3, stride=2),
            nn.Conv2d(64, 192, kernel_size=5, padding=2),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=3, stride=2),
            nn.Conv2d(192, 384, kernel_size=3, padding=1),
            nn.ReLU(inplace=True),
            nn.Conv2d(384, 256, kernel_size=3, padding=1),
            nn.ReLU(inplace=True),
            nn.Conv2d(256, 256, kernel_size=3, padding=1),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=3, stride=2)
        )

        self.classifier = nn.Sequential(
            nn.Dropout(),
            nn.Linear(256 * 6 * 6, 4096),
            nn.ReLU(inplace=True),
            nn.Dropout(),
            nn.Linear(4096, 4096),
            nn.ReLU(inplace=True),
            nn.Linear(4096, num_classes),
        )

        self.load_weights()

    # we modified this section
    def forward(self, x):

```



```
x = self.features(x)
x = x.view(x.size(0), 256 * 6 * 6)
```

Extract the values of the activations of AlexNet on the face images in a particular layer

Input images size: 227*227

We will extract activations of the network during the forward pass.

Then, we use part 8 model to train a fully-connected neural network by taking the activations as input.

Get the activation values as input to part8 code as follows:

Listing 9: How to get activation values

5

```
# get train and test activations using forward
x_train = Variable(torch.from_numpy(train_set[:]), requires_grad=False).type(torch.FloatTensor)
train_activation = np.vstack((train_activation, model.forward(x_train).data.numpy()))
x_test = Variable(torch.from_numpy(train_set[:]), requires_grad=False).type(torch.FloatTensor)
test_activation = np.vstack((test_activation, model.forward(x_test).data.numpy()))

train_using_mini_batches(act, train_activation, test_activation, "part10", 227)
```

Performance:

The performance dramatically improved compared to part8.

Using myAlexNet model, we were able to obtain a final performance of 95.83

The following is the performance:

```
('epoch: ', 0, 'train performance:', 55.21172638436482, 'test performance: ', 45.83333333333333)
('epoch: ', 50, 'train performance:', 100.0, 'test performance: ', 95.83333333333334)
('epoch: ', 100, 'train performance:', 100.0, 'test performance: ', 95.83333333333334)
('epoch: ', 150, 'train performance:', 100.0, 'test performance: ', 95.83333333333334)
('epoch: ', 200, 'train performance:', 100.0, 'test performance: ', 95.83333333333334)
('epoch: ', 250, 'train performance:', 100.0, 'test performance: ', 95.83333333333334)
('epoch: ', 300, 'train performance:', 100.0, 'test performance: ', 95.83333333333334)
('epoch: ', 350, 'train performance:', 100.0, 'test performance: ', 95.83333333333334)
('epoch: ', 400, 'train performance:', 100.0, 'test performance: ', 95.83333333333334)
('epoch: ', 450, 'train performance:', 100.0, 'test performance: ', 95.83333333333334)
('epoch: ', 499, 'train performance:', 100.0, 'test performance: ', 95.83333333333334)
```

The performance improved by more than 30% compared to part 8.

Run the part10 code using: python myalexnet.py